



The Trainer's Friend
I N C O R P O R A T E D

An Introduction to Unicode

- ◆ **Unicode Concepts and Terminology**
- ◆ **Unicode Mappings**
- ◆ **Appendix: UTF-32 Character Assignment Ranges**
- ◆ **Appendix: UTF-32 <-> UTF-16 <-> UTF-8 sample mappings**

by Steve Comstock

The Trainer's Friend, Inc.
<http://www.trainersfriend.com>
303-393-8716
steve@trainersfriend.com

v1.7

The following terms that may appear in these course materials are trademarks or registered trademarks:

Trademarks of the International Business Machines Corporation:

AD/Cycle, AIX, AIX/ESA, Application System/400, AS/400, BookManager, CICS, CICS/ESA, COBOL/370, COBOL for MVS and VM, COBOL for OS/390 & VM, Common User Access, CORBA, CUA, DATABASE 2, DB2, DB2 Universal Database, DFSMS, DFSMSds, DFSORT, DOS/VSE, Enterprise System/3090, ES/3090, 3090, ESA/370, ESA/390, Hiperbatch, Hiperspace, IBM, IBMLink, IMS, IMS/ESA, Language Environment, MQSeries, MVS, MVS/ESA, MVS/XA, MVS/DFP, NetView, NetView/PC, Object Management Group, Operating System/400, OS/400, PR/SM, OpenEdition MVS, Operating System/2, OS/2, OS/390, OS/390 UNIX, OS/400, QMF, RACF, RS/6000, SOMobjects, SOMobjects Application Class Library, System/370, System/390, Systems Application Architecture, SAA, System Object Model, TSO, VisualAge, VisualLift, VTAM, VM/XA, VM/XA SP, WebSphere, z/OS, z/VM, z/Architecture, zSeries

Trademarks of Microsoft Corp.: Microsoft, Windows, Windows NT, Visual Basic, Microsoft Access, MS-DOS, Windows XP

Trademarks of Micro Focus Corp.: Micro Focus

Trademarks of America Online, Inc.: America Online

Trademarks of Quercus Systems: Personal REXX, REXXTERM

Trademark of Chicago-Soft, Ltd: MVS/QuickRef

Trademark of Crystal Computer Services: Crystal Reports

Trademark of Phoenix Software International: (E)JES

Registered Trademarks of Institute of Electrical and Electronic Engineers: IEEE, POSIX

Registered Trademarks of Corel Corporation: Corel, CorelDRAW, Corel VENTURA

Registered Trademark of Oracle Corporation: Oracle

Registered Trademark of The Open Group: UNIX

Trademark of Sun Microsystems, Inc.: Java

Registered Trademark of Linus Torvalds: LINUX

Registered Trademark of Unicode, Inc.: Unicode

Trademarks held on behalf of World Wide Web Consortium: W3C, XHTML, XSL, WebFonts

Section Preview

☐ Introduction to Unicode

- ◆ Characters
- ◆ Characters, Glyphs, and Fonts
- ◆ Coding Schemes
- ◆ Code pages
- ◆ Standards
- ◆ Unicode

☐ Processing Unicode Characters

Characters

- ❑ A character is an abstract concept that has evolved along with our understanding of language and information

- ❑ Initially, when most of us think of characters, we think of a particular character set: alphabetic letters, numeric characters, special characters, and so on
 - ◆ But subtleties start to appear when you consider issues of other languages, different fonts, and the need for special meaning characters
 - ✗ For example, you need characters to control a communications session when transmitting data

 - ✗ Obviously, you cannot use characters from the set of data you are sending to also be control characters: the communications process could not distinguish between data and the control functions

Characters, Glyphs, and Fonts

☐ In computer terms, a character is a grouping of bits (binary ones and zeros) in packages of 8: one or more bytes

☐ There are two broad classes of characters: data characters and control characters

♦ Although one could make a case that data characters are just control characters whose function is to display a glyph

✕ A glyph is the visible representation of a character

♦ Consider the [data] character called "upper case A"; the following are various glyphs that represent that character:

✕ A - Arial

✕ A - Times New Roman

✕ A - Courier new

✕ A - Garamond

✕ A - Bodoni

✕ A - Park Avenue

✕ A - FlamencoD

♦ And so on; notice the concept of a font sneaking in here: a font is a set of glyphs used to represent a collection of characters [usually in a similar style]

Coding Schemes

- ☐ But an upper case A is an upper case A, regardless of the glyph used to represent it
- ☐ In computers, we assign characters to bit patterns (and vice versa), and a "character" is an abstract thing, independent of any glyph
- ☐ The rules we use to make these assignments between characters and bit patterns are called coding schemes, and there are many in use today, for historical reasons
- ☐ There are three coding schemes most people in the IS industry need to be cognizant of:
 - ◆ EBCDIC - Extended Binary Coded Decimal Interchange Code; used by IBM mainframes and AS/400 machines
 - ◆ ASCII - American Standard Code for Information Interchange; used by almost all other hardware
 - ◆ Unicode - gaining wide acceptance in use by software
- ☐ There are other coding schemes available, but from a practical point of view, we can get the vast majority, if not all, of our work done if we are aware of these coding schemes

Codepages

- ❑ Even awareness of coding schemes is not quite enough to get us all we need for practical use
- ❑ Again, for historical and cultural reasons, many coding schemes have several variations, each slightly different than the others
 - ◆ For example, in some environments you have need of a symbol like Å, but in other environments, users are not even aware of this character
 - ◆ So computer designers introduced the concept of a codepage, which is a variation of a coding scheme
- ❑ After all, in 8 bits you only have 256 possible patterns
 - ◆ You can run out of available characters pretty quick if you allow all those strange foreign, mathematical, scientific, engineering, currency, and other symbols
- ❑ The solution was to use codepages (also spelled as two words: "code page" or "code pages")
 - ◆ Users could set codepages for different environments
 - ◆ Although you cannot mix codepages in a single environment: at any point in time your 256 bit patterns map to exactly one set of characters

Codepages, continued

❑ When a data character arrives in a computer from a magnetic tape, diskette, CD-ROM, network transmission, or so on, the computer just stores the character as it comes in, no judgements being made

❑ But consider what happens when a character comes in from a keyboard:

- ◆ The user presses a key with a glyph on it of, say, an upper case A
- ◆ The keyboard electronics assign a bit pattern to the character and transmit it to the computer, where it is received as part of an I/O program
- ◆ This I/O program may reassign the bit pattern before it is stored, depending on the current codepage:

keyboard bit pattern → **codepage** → stored bit pattern

❑ Similarly, when a character is sent by the computer to a printer or display unit, that output device has a codepage mapping followed by a font mapping to decide how to display the character on the device

stored bit pattern → **codepage** → character → **font** → glyph

Standards

- ❑ To ensure consistency and clarity, a number of standards bodies have been created to develop and enhance standards for a variety of areas, including IS; these bodies include:

- ◆ ISO - the International Organization for Standardization
- ◆ ANSI - the American National Standards Institute, is the US member of ISO
- ◆ The ISO has a standard called ISCII which is very close to the ASCII character encoding standard

- ❑ Ultimately, one wants a single code page, a single, universal, encoding scheme for all characters

- ❑ From the perspective of international communications, one needs an encoding scheme that is

- ◆ Universal - covers all characters needed in all likely situations
- ◆ Efficient - avoids escape character sequences for special encoding, for example
- ◆ Unambiguous - every character has one and only one bit pattern mapping

Unicode

- ☐ **Unicode is an encoding scheme developed by the Unicode Consortium (incorporated under the name Unicode, Inc. in 1991) and the ISO**

- ☐ **The Unicode Consortium is backed by most of the major players in the IS game, including these (and many more):**
 - ◆ **Adobe Systems**
 - ◆ **Apple Computer**
 - ◆ **Compaq Computer**
 - ◆ **Ericsson Mobile Communications**
 - ◆ **Hewlett-Packard**
 - ◆ **IBM**
 - ◆ **Microsoft**
 - ◆ **NCR**
 - ◆ **Netscape**
 - ◆ **Oracle**
 - ◆ **PeopleSoft**
 - ◆ **Quark**
 - ◆ **SAP**
 - ◆ **SAS Institute**
 - ◆ **SHARE**
 - ◆ **Software AG**
 - ◆ **Sun Microsystems**
 - ◆ **Sybase**
 - ◆ **Unisys**

Unicode, continued

- ❑ In 1992, the Unicode Consortium and ISO agreed to merge their character encoding standard, so the character sets map exactly
 - ◆ In addition to assigning names and bit-pattern mappings to characters, in conjunction with the ISO, the Unicode standard also provides implementation algorithms, properties, and semantic information

- ❑ The basic, original premise, was to use 16-bits for every character
 - ◆ This allows for 64K unique patterns (65,536)
 - ◆ Maintaining compatibility with as many already existing standards as possible

- ❑ By the year 2000, however, it was clear that more character space was needed
 - ◆ In May of 2001, 44,946 new characters were added (mostly CJK (Chinese, Japanese, Korean) characters, along with some historic scripts and several sets of symbols)
 - ✗ As of Unicode standard 3.1 there were 94,140 characters included
 - ◆ As of Unicode standard 4.0 (June, 2003), there are 96,382 characters in the standard, and for 4.1 (March, 2005) the count is now 97,655; for 5.0 (July, 2006), there are 99,024

Unicode, continued

❑ There are three alternative ways of representing Unicode characters, including:

- ◆ UTF-16 — the basic 16-bit encoding scheme: two bytes used for every Unicode character

- ✗ But version 3.0 of the Unicode standard introduced a concept called surrogate pairs that allows some Unicode characters to be represented by a pair of two-byte values

- ◆ UTF-8 — an algorithm for converting Unicode characters to a string of characters that are one, two, three, or four bytes in length, and back

- ◆ UTF-32 — a 32-bit encoding, the basis for the ultimate character encoding, allowing for 1,114,112 character assignments (note: the leftmost 11 of the 32 bits must be all binary zeros)

- ✗ This encoding was made an official part of the Unicode standard in version 3.1 in May, 2001

❑ UTF stands for Unicode Transformation Format

❑ Here are some pointers to Web sites for more information about Unicode:

- ◆ Unicode home page: <http://www.unicode.org>

- ◆ IBM: <http://www-106.ibm.com/developerworks/unicode/>

Unicode, continued

☐ So why do we care about this on the mainframe?

- ◆ **IBM is trying to position mainframes as the ultimate server for intranets and the Internet / World Wide Web**

- ✗ Web pages are generally coded in HTML (HyperText Markup Language) or XHTML (eXtensible HyperText Markup Language)

- HTML 4 and all versions of XHTML require support for Unicode

- ◆ **XML (eXtensible Markup Language) is becoming one of the premier data exchange formats - requires Unicode support**

- ◆ **At some point in time, ("not too far down the road" to quote one of the z/OS developers) z/OS will require the Unicode support functions be installed**

- ◆ **DB2 can store / access Unicode data in CHAR, VARCHAR, and CLOB data types**

- ✗ In Version 8, the DB2 catalog is stored in UTF-8

- ◆ **Many databases and programming languages on UNIX and Windows machines support Unicode**

- ◆ **Current mainframe compilers (COBOL, PL/I, C, C++) all support Unicode**

☐ **Unicode can provide, eventually, the ability to have a single codepage yet support all languages simultaneously**

Unicode, concluded

❑ Although there are people who are against Unicode (and even some competing standards), Unicode seems to be the way of the future

- ◆ Enabling single encoding and data interchange across platforms and simultaneous multiple language support on screens and reports

❑ Also note that z-series machines have a number of instructions that work with Unicode, in the UTF-16 format (PKU, UNPKU, CLCLU, MVCLU, TROO, TROT, TRTO, TRTT)

- ◆ But UTF-8 seems to be the most widely used format on the Web and, probably, in XML that's not even used on the Web
- ◆ Instructions to convert between UTF-16 and UTF-8 have been available since machines introduced in 1999 (CUTFU, CUUTF)
- ◆ In 2004, instructions were added to convert between:

X UTF-8 <--> UTF-16 (new names for old instructions)

X UTF-8 <--> UTF-32 (new instructions)

X UTF-16 <--> UTF-32 (new instructions)

Section Preview

☐ Processing Unicode Characters

- ◆ Unicode Representations
- ◆ UTF-32: Unicode Scalar Values
- ◆ UTF-16
- ◆ Surrogate Pairs
- ◆ UTF-16 -> Unicode Scalar Value
- ◆ UTF-32 -> UTF-16
- ◆ UTF-8
- ◆ UTF-32 -> UTF-8
- ◆ . UTF-8 -> UTF-32
- ◆ . Other Mappings
- ◆ . Unicode - Conclusion

Unicode Representations

- ❑ **This section is a technical discussion of how Unicode characters are stored using the three formats: UTF-8, UTF-16, and UTF-32**

- ◆ **According to the standard, these are considered equally valid**

- ✗ In the sense that every Unicode character may be represented in any of these formats, and the mapping between formats is well-defined

- ◆ **To work with Unicode data, one has to know which encoding format has been used**

- ✗ This may be supplied external to the data itself, as in an HTTP header or HTML META statement, for example

- ✗ In some cases, the program processing a string may be able to examine the string and deduce the format being used (but this is not preferred)

UTF-32: Unicode Scalar Values

- ❑ Every Unicode character is assigned an integer value, the Unicode scalar value

- ❑ UTF-32 is the set of Unicode Scalar Values

- ◆ This is also sometimes called UCS-4, meaning the Universal Character Set as 4-bytes per character

- ❑ The possible range of values in the Unicode scalar set is `x'00000000' - x'0010FFFF'` or, in binary, the maximum allowed value is

0000 0000 0001 0000 1111 1111 1111 1111

- ◆ 21 bits; in decimal, the values range from 0 to 1,114,111
- ◆ Every Unicode character is assigned a number in this range, a point along the string of integers in this range (this is sometimes called a code point)
 - ✗ Although not every number in this range is assigned to a Unicode character
 - ✗ Also, a subset of this range is reserved for surrogate pairs...

UTF-16

- ❑ UTF-16 was the beginning point of Unicode character assignments
- ❑ Initially, each UTF-16 character was a single two-byte unit
 - ◆ But when surrogates needed to be introduced, to accomodate a larger character set, some characters became represented by a single two-byte unit, others by a pair of two-byte units
- ❑ When a processing program such as a browser or editor is working with UTF-16 data, it assumes each two-byte unit represents a character
 - ◆ Except that certain values are reserved to represent surrogate pairs: situations where it takes two two-byte units to represent a single character
- ❑ The theoretical range of Unicode scalar values is `x'0000 0000' - x'0010 FFFF'`
 - ◆ For Unicode scalar values greater than or equal to `x'0001 0000'`, surrogate pairs are used
 - ◆ Values in the range `x'0000 D800' - x'0000 DFFF'` are reserved for use in surrogate pairs

Surrogate Pairs

❑ How to recognize when a two-byte unit begins a surrogate pair?

- ◆ If a UTF-16 unit has a value in the range `x'D800' - x'DBFF'`, that unit is a high surrogate and you need to combine that two-byte unit and the next, which must be a low surrogate, to determine the actual character that is being represented

- ◆ Low surrogates are in the range `x'DC00' - x'DFFF'`

✗ It is an error for a low surrogate not to be preceded by a high surrogate, and for a high surrogate not to be followed by a low surrogate

In binary:

high surrogates: 1101 1000 0000 0000 - 1101 1011 1111 1111

low surrogates: 1101 1100 0000 0000 - 1101 1111 1111 1111

In decimal:

high surrogates: 55,296 - 56,319

low surrogates: 56,320 - 57,343

Notes

- ◆ Each surrogate range contains 1,024 values, so the possible number of surrogate pair values is 1,024 x 1,024 or 1,048,576
- ◆ The vast majority of characters do not require surrogate pairs

UTF-16 -> Unicode Scalar Value

□ The algorithm to convert from a UTF-16 Unicode character to the Unicode scalar value (in other words, UTF-16 -> UTF-32) is this:

♦ If a two-byte unit is not a surrogate value, the Unicode scalar value is the two-byte value itself

X So if the two-byte unit is in the range $x'0000' - x'D7FF'$ or $x'E000' - x'FFFF'$, the Unicode scalar value is that value (or, equivalently, $x'0000\ 0000' - x'0000\ D7FF'$ and $x'0000\ E000' - x'0000\ FFFF'$)

♦ If a two-byte unit is a surrogate value, the character is composed of two two-byte units, so calculate the Unicode scalar value as the sum of

X (The high surrogate - $x'D800'$) * $x'0400'$

X (The low surrogate - $x'DC00'$)

X $x'0001\ 0000'$

□ We examine this algorithm more carefully ...

UTF-16 -> Unicode Scalar Value, continued

Notes

- ◆ The first calculation provides the displacement into the high surrogate range (resulting in a number in the range x'0000'- x'03FF' or, in decimal, 0 - 1023 or, in binary: b'0000 0000 0000 0000'- b'0000 0111 1111 1111')
- ◆ Multiplying by x'0400' (decimal 1024) effectively shifts the value to the left 10 bits, producing numbers in the range x'0000 0000'- x'000 FFC00' with the last 10 bits all zeros

0000 00xx xxxx xxxx

0000 0000 xxxx xxxx xx00 0000 0000

- ◆ The second value is the displacement into the low surrogate range (also resulting in a number in the range x'0000'- x'03FF' or, in decimal, 0 - 1023)
- ◆ Adding the two numbers (inserting leading zeros in the first to make them the same length) and the x'1 0000':

```
0000 0000 0000 xxxx xxxx xx00 0000 0000
0000 0000 0000 0000 0000 00yy yyyY yyyY
+ 0000 0000 0000 0001 0000 0000 0000 0000
```

- ◆ Adding the x'0001 0000' ensures the resulting Unicode scalar values are in the range x'0001 0000'- x'0010 FFFF'

UTF-16 -> Unicode Scalar Value, continued

- We can look at Unicode scalar assignments this way:

x'0000 0000'	-	x'0000 D7FF'	basic 16-bit codes
x'0000 D800'	-	x'0000 DFFF'	surrogate values (assigned, but not to characters)
x'0000 E000'	-	x'0000 FFFF'	basic 16-bit codes
x'0001 0000'	-	x'0010 FFFF'	computed from surrogate pairs

- These ranges may be further subdivided for study but these subsets are not of interest in this paper

- ◆ However, the next level of detail is presented in the first Appendix to this document
- ◆ The second Appendix lists specific mapping values between UTF-32, UTF-16, and UTF-8

- At one point in the cycle of development, there was a Unicode encoding called UCS-2 (Universal Character Set as 2-bytes per character)

- ◆ UCS-2 is, essentially, UTF-16 without support for surrogate pairs

✗ UCS-2 is currently supplanted by UTF-16

UTF-32 -> UTF-16

❑ On the other hand, given a UTF-32 character, how do you represent it in UTF-16?

♦ This algorithm, of course, reverses the steps before:

✗ For characters less than $x'0001\ 0000'$, the UTF-16 representation is simply the rightmost 16 bits

✗ For characters in the range $x'0001\ 0000'$ to $x'0010\ FFFF'$ we need to build the high surrogate (HS) and low surrogate (LS) two-byte patterns, as follows...

➤ Subtract $x'0001\ 0000'$; this gives a value in the range $x'0000\ 0000'$ to $x'000F\ FFFF'$, call this value **char**

➤ $HS = x'D800' + (char / x'400')$

➤ $LS = x'DC00' + (char \% x'400')$

✗ In other words, consider the 20 rightmost bits of **char** to be designated this way:

xxxx xxxx xx|yy yyyy yyyy

✗ then the HS is $x'D800'$ plus the leftmost 10 bits (the x's) and the LS is $x'DC00'$ plus the rightmost 10 bits (the y's)

UTF-8

- ☐ One of the major reasons Unicode has been successful is a deliberate decision to encompass as many existing standards as possible

- ☐ When it came to the 7-bit ASCII / ISCII standard, Unicode allowed an 8-bit representation for characters with binary code points from b'0000 0000' to b'0111 1111'
 - ◆ In other words, the first 127 Unicode characters are ASCII!

- ☐ UTF-8 allows you to represent any Unicode character as a string of one, two, three, or four bytes!


- ☐ Conversely, when a processor knows it is dealing with UTF-8, it starts by looking at one byte and the value will imply whether it needs to use one, two, three, or four bytes to construct the UTF-32 value (Unicode scalar value)

UTF-32 -> UTF-8

❑ The mapping works like this


- ◆ Unicode scalar values in the range 00000000-0000007F map to single byte values 00-7F
- ◆ Unicode scalar values in the range 00000080-000007FF map to two byte values, where the first byte is in the range C2-DF and the second byte is in the range 80-BF

✕ Specifically looking at the bit patterns:



- ◆ Unicode scalar values in the range 00000800-00000FFF map to three byte values, where the first byte is E0, the second byte is in the range A0-BF, and the third byte is in the range 80-BF

✕ Specifically looking at the bit patterns:



UTF-32 -> UTF-8, continued

❑ The mapping works like this, continued


- ◆ Unicode scalar values in the range 00001000-0000FFFF map to three byte values, where the first byte is in the range E1-EF, the second byte is in the range 80-BF, and the third byte is in the range 80-BF

✕ Specifically looking at the bit patterns:



- ◆ Unicode scalar values in the range 00010000-0003FFFF map to four byte values, where the first byte is F0, the second byte is in the range 90-BF, the third byte is in the range 80-BF, and the fourth byte is in the range 80-BF

✕ Specifically looking at the bit patterns:




UTF-32 -> UTF-8, continued

❑ The mapping works like this, continued


- ◆ Unicode scalar values in the range 00040000-000FFFFF map to four byte values, where the first byte is in the range F1-F3, the second byte is in the range 80-BF, the third byte is in the range 80-BF, and the fourth byte is in the range 80-BF

✕ Specifically looking at the bit patterns:



- ◆ Unicode scalar values in the range 00100000-0010FFFF map to four byte values, where the first byte is F4, the second byte is in the range 80-BF, the third byte is in the range 80-BF, and the fourth byte is in the range 80-BF

✕ Specifically looking at the bit patterns:



UTF-8 -> UTF-32

- ❑ Obviously, the reverse mapping works based on the value a processor finds in the first byte of a UTF-8 string:
 - ◆ If a byte is in the range 00-7F, it constructs the UTF-32 output by prefixing hex 000000
 - ◆ If a byte is in the range C2-DF, it knows to take that byte and the next to build the UTF-32 value
 - ◆ If a byte is E0-EF, it knows to take that byte and the next two bytes to build the UTF-32 value
 - ◆ If a byte is F0-F4, it knows to take that byte and the next three to build the UTF-32 scalar value and then to map that to the surrogate pair
 - ◆ Any other values in the first byte indicate an error
- ❑ The mechanics should be reasonably apparent from the preceding pages, and details are left as a task for the interested reader to do on your own
 - ◆ The second Appendix to this document shows many code point mappings between UTF-32, UTF-16, and UTF-8

Other Mappings

- ☐ We have not discussed these mappings
 - ◆ UTF-8 -> UTF-16
 - ◆ UTF-16 -> UTF-8

- ☐ You can accomplish these mappings by using UTF-32 as an intermediate format then using the existing UTF-32 <-> UTF-16 mappings
 - ◆ Or it's possible to devise direct mappings based on what we've discussed

- ☐ On IBM z/Architecture machines, there are instructions that convert between Unicode formats:
 - ◆ CU12 - from UTF-8 to UTF-16 (also known as CUTFU)
 - ◆ CU21 - from UTF-16 to UTF-8 (also known as CUUTF)
 - ◆ CU14 - from UTF-8 to UTF-32
 - ◆ CU41 - from UTF-32 to UTF-8
 - ◆ CU24 - from UTF-16 to UTF-32
 - ◆ CU42 - from UTF-32 to UTF-16

- ☐ Our interest in this paper has been simply in demonstrating the interrelationships between the three Unicode formats and we feel that has been explored thoroughly enough at this point

Endian-ness

- ❑ One issue we didn't raise that needs to be raised in certain circumstances: UTF-16 and UTF-32 each have variants depending on the order the bytes are physically stored

- ◆ In mainframe (and many other systems) these characters are stored in Big Endian (BE) format: most significant byte first

✗ Officially called UTF-16BE and UTF-32BE

- ◆ In other systems, characters are stored in Little Endian (LE) format: least significant byte first

✗ Officially designated UTF-16LE and UTF-32LE

- ❑ When passing UTF-16 and UTF-32 strings, one has to specify the "endian-ness" of the strings, in headers or by inserting a hex value called a Byte Order Mark (BOM) at the front of the data

- ◆ In UTF-16, an initial byte sequence of `x'FEFF'` indicates big endian encoding, while `x'FFFE'` indicates little endian encoding
- ◆ In UTF-32, an initial byte sequence of `x'0000 FEFF'` indicates big endian encoding, while `x'FFFE 0000'` indicates little endian order
- ◆ In both cases, any BOM is not considered part of the data, and the absence of any BOM value implies big endian encoding

Unicode - Conclusion

- ❑ **Unicode is here, is well-defined, and is gaining wide acceptance on a variety of hardware and software platforms**

- ❑ **Although new characters and additional refinements continue to be made, the Unicode Consortium strives to make each new version backward compatible**
 - ◆ **In other words, current algorithms and mappings will carry forward**

- ❑ **There is plenty of work to do for those interested in exploring the internationalization of the Web and of computer work in general**
 - ◆ **The ultimate goal is effective communication between men and women of all languages and cultures on the planet**

Section Preview

☐ Appendix A

◆ UTF-32 character allocation ranges

UTF-32 Character Allocation Ranges

❑ Every Unicode character is assigned a scalar value: an integer

◆ This is a 21-bit number in the range:

binary
000000000000000000000 - 1 0000 1111 1111 1111 1111

or

hex
00 00 00 - 10 FF FF

decimal
0 - 1,114,111

UTF-32 Assignments, continued

- ❑ UTF-32 is really just the 21-bit numbers, right justified, padded on the left with 11 binary zeros (so only six relevant hex digits)

◆ General allocation is this (not all gaps and details are shown):

00 00 00 - 00 00 7F	7-bit ASCII	
00 00 80 - 00 00 FF	Controls and Latin-1 Supplement	
00 01 00 - 00 01 7F	Latin Extended-A	
00 01 80 - 00 02 4F	Latin Extended-B	
00 02 50 - 00 02 AF	IPA extensions (some Latin and Greek)	
00 02 B0 - 00 02 FF	Spacing modifier letters	
00 03 00 - 00 03 6F	Combining diacritical marks	
00 03 70 - 00 03 FF	Greek and Coptic	
00 04 00 - 00 04 FF	Cyrillic	
00 05 00 - 00 05 2F	Cyrillic supplementary	
00 05 30 - 00 05 8F	Armenian	
00 05 90 - 00 05 FF	Hebrew	
00 06 00 - 00 06 FF	Arabic	
00 07 00 - 00 07 4F	Syriac	
00 07 50 - 00 07 7F	Thaana	(note gap here)
00 09 00 - 00 09 7F	Devanagari	
00 09 80 - 00 09 FF	Bengali	
00 0A 00 - 00 0A 7F	Gurmukhi	
00 0A 80 - 00 0A FF	Gujarati	
00 0B 00 - 00 0B 7F	Oriya	
00 0B 80 - 00 0B FF	Tamil	
00 0C 00 - 00 0C 7F	Telugu	
00 0C 80 - 00 0C FF	Kannada	
00 0D 00 - 00 0D 7F	Malayalam	
00 0D 80 - 00 0D FF	Sinhala	
00 0E 00 - 00 0E 7F	Thai	
00 0E 80 - 00 0E FF	Lao	
00 0F 00 - 00 0F FF	Tibetan	
00 10 00 - 00 10 9F	Myanmar	
00 10 A0 - 00 10 FF	Georgian	
00 11 00 - 00 11 FF	Hangul Jamo	

UTF-32 Assignments, continued

☐ UTF-32 scalar values allocation, continued:

00 12 00 - 00 13 7F	Ethiopic	
00 13 A0 - 00 13 FF	Cherokee	
00 14 00 - 00 16 7F	Unified Canadian Aboriginal symbols	
00 16 80 - 00 16 9F	Ogham	
00 16 A0 - 00 16 FF	Runic	
00 17 00 - 00 17 1F	Tagalog	
00 17 20 - 00 17 3F	Hanunoo	
00 17 40 - 00 17 5F	Buhid	
00 17 60 - 00 17 7F	Tagbanwa	
00 17 80 - 00 17 FF	Khmer	
00 18 00 - 00 18 AF	Mongolian	(note gap here)
00 19 00 - 00 19 4F	Limbu	
00 19 50 - 00 19 7F	Tai Le	(note gap here)
00 19 E0 - 00 19 FF	Khmer symbols	(note gap here)
00 1D 00 - 00 1D 7F	Phonetic extensions	(note gap here)
00 1E 00 - 00 1E FF	Latin extended additional	
00 1F 00 - 00 1F FF	Greek extended	
00 20 00 - 00 20 6F	General punctuation	
00 20 70 - 00 20 9F	Superscripts and subscripts	
00 20 A0 - 00 20 CF	Currency symbols	
00 20 D0 - 00 20 FF	Combining diacritical marks for symbols	
00 21 00 - 00 21 4F	Letterlike symbols	
00 21 50 - 00 21 8F	Number forms	
00 21 90 - 00 21 FF	Arrows	
00 22 00 - 00 22 FF	Mathematical operators	
00 23 00 - 00 23 FF	Miscellaneous technical	
00 24 00 - 00 24 3F	Control pictures	
00 24 40 - 00 24 5F	Optical character recognition	
00 24 60 - 00 24 FF	Enclosed alphanumerics	
00 25 00 - 00 25 7F	Box drawing	
00 25 80 - 00 25 9F	Block elements	
00 25 A0 - 00 25 FF	Geometric shapes	
00 26 00 - 00 26 FF	Miscellaneous symbols	

UTF-32 Assignments, continued

☐ UTF-32 scalar values allocation, continued:

00 27 00 - 00 27 BF	Dingbats
00 27 C0 - 00 27 EF	Miscellaneous mathematical symbols-A
00 27 F0 - 00 27 FF	Supplemental arrows-A
00 28 00 - 00 28 FF	Braille patterns
00 20 00 - 00 29 7F	Supplemental arrows-B
00 29 80 - 00 29 FF	Miscellaneous mathematical symbols-B
00 2A 00 - 00 2A FF	Supplemental mathematical operators
00 2B 00 - 00 2B FF	Miscellaneous symbols and arrows
00 2C 00 - 00 2E 7F	unassigned
00 2E 80 - 00 2E FF	CJK radicals supplement
00 2F 00 - 00 2F DF	Kangxi radicals (note gap here)
00 2F F0 - 00 2F FF	Ideographic description characters
00 30 00 - 00 30 3F	CJK symbols and punctuation
00 30 40 - 00 30 9F	Hiragana
00 30 A0 = 00 30 F0	Katakana
00 31 00 - 00 31 2F	Bopomofo
00 31 30 - 00 31 8F	Hangul compatibility Jamo
00 31 90 - 00 31 9F	Kanbun
00 31 A0 - 00 31 BF	Bopomofo extended (note gap here)
00 31 F0 - 00 31 FF	Katakana phonetic extensions
00 32 00 - 00 32 FF	Enclosed CJK letters and months
00 33 00 - 00 33 FF	CJK compatibility
00 34 00 - 00 4D BF	CJK unified ideographs extension A
00 4D C0 - 00 4D FF	Yijing Hexagram symbols
00 4E 00 - 00 9F AF	CJK unified ideographs
00 A0 00 - 00 A4 8F	Yi syllables
00 A4 90 - 00 A4 CF	Yi radicals (note gap here)
00 AC 00 - 00 D7 AF	Hangul syllables (note gap here)
00 D8 00 - 00 DB FF	high surrogates
00 DC 00 - 00 DF FF	low surrogates (note gap here)
00 F9 00 - 00 FA FF	CJK compatibility ideographs
00 FB 00 - 00 FB 4F	Alphabetic presentation forms
00 FB 50 - 00 FD FF	Arabic presentation forms-A

UTF-32 Assignments, continued

☐ UTF-32 scalar values allocation, continued:

00 FE 00 - 00 FE 0F	Variation selectors	
00 FE 20 - 00 FE FF	Combining half marks	
00 FE 30 - 00 FE 4F	CJK compatibility forms	
00 FE 50 - 00 FE 6F	Small form variants	
00 FE 70 - 00 FE FF	Arabic presentation forms-B	
00 FF 00 - 00 FF EF	Halfwidths and fullwidth forms	
00 FF F0 - 00 FF FF	Specials	
<hr/>		
01 00 00 - 01 00 7F	Linear B syllabary	
01 00 80 - 01 00 FF	Linear B ideograms	
01 01 00 - 01 01 3F	Aegean numbers	(note gap here)
01 03 00 - 01 03 2F	Old Italic	
01 03 30 - 01 03 4F	Gothic	(note gap here)
01 03 80 - 01 03 9F	Ugaritic	(note gap here)
01 04 00 - 01 04 4F	Deseret	
01 04 50 - 01 04 7F	Shavian	
01 04 80 - 01 04 AF	Osmanya	(note gap here)
01 08 00 - 01 08 3F	Cypriot syllabary	(note gap here)
01 D0 00 - 01 D0 FF	Byzantine musical symbols	
01 D1 00 - 01 D1 FF	Musical symbols	(note gap here)
01 D3 00 - 01 D3 5F	Tai Xuan Jing symbols	(note gap here)
01 D4 00 - 01 D7 FF	Mathematical alphanumeric symbols	
01 D8 00 - 01 FF FF	unassigned	
02 00 00 - 02 A6 DF	CJK unified ideographs extension B	
02 A6 E0 - 02 F7 FF	unassigned	
02 F8 00 - 02 FA 1F	CJK compatibility ideographs supplement	
02 FA 20 - 0D FF FF	unassigned	
0E 00 00 - 0E 00 7F	Tags	(note gap here)
0E 01 00 - 0E 01 EF	Variant selectors supplement	
0E 01 F0 - 0E FF FF	unassigned	
0F 00 00 - 0F FF FD	Private use area	
0F FF FE - 0F FF FF	non-characters	
10 00 00 - 10 FF FD	Private use area	
10 FF FE - 10 FF FF	non-characters	

This page intentionally left almost blank.

Section Preview

☐ Appendix B

- ◆ UTF-32 <-> UTF-16 <-> UTF-8 sample mappings

Mappings

UTF-32	UTF-16	UTF-8
00 00 00	00 00	00
00 00 01	00 01	01
.		
.		
.		
00 00 7E	00 7E	7E
00 00 7F	00 7F	7F
00 00 80	00 80	C2 80
00 00 81	00 81	C2 81
.		
.		
.		
00 00 BE	00 BE	C2 BE
00 00 BF	00 BF	C2 BF
00 00 C0	00 C0	C3 80
00 00 C1	00 C1	C3 81
.		
.		
.		
00 00 FE	00 FE	C3 BE
00 00 FF	00 FF	C3 BF
00 01 00	01 00	C4 80
00 01 01	01 01	C4 81
.		
.		
.		
00 01 3E	01 3E	C4 BE
00 01 3F	01 3F	C4 BF
00 01 40	01 40	C5 80
00 01 41	01 41	C5 81
.		
.		
.		

Mappings, 2

UTF-32	UTF-16	UTF-8
00 01 7E	01 7E	C5 BE
00 01 7F	01 7F	C5 BF
00 01 80	01 80	C6 80
00 01 81	01 81	C6 81
.		
.		
.		
00 01 BE	01 BE	C6 BE
00 01 BF	01 BF	C6 BF
00 01 C0	01 C0	C7 80
00 01 C1	01 C1	C7 81
.		
.		
.		
00 01 FE	01 FE	C7 BE
00 01 FF	01 FF	C7 BF
00 02 00	02 00	C8 80
00 02 01	02 01	C8 81
.		
.		
.		
00 02 3E	02 3E	C8 BE
00 02 3F	02 3F	C8 BF
00 02 40	02 40	C9 80
00 02 41	02 41	C9 81
.		
.		
.		
00 02 7E	02 BE	C9 BE
00 02 7F	02 BF	C9 BF
00 02 80	02 00	CA 80
00 02 81	02 01	CA 81
.		
.		
.		

Mappings, 3

UTF-32	UTF-16	UTF-8
00 02 BE	02 BE	CA BE
00 02 BF	02 BF	CA BF
00 02 C0	02 C0	CB 80
00 02 C1	02 C1	CB 81
.		
.		
.		
00 02 FE	02 FE	CB BE
00 02 FF	02 FF	CB BF
00 03 00	03 00	CC 80
00 03 01	03 01	CC 81
.		
.		
.		
00 03 3E	03 3E	CC BE
00 03 3F	03 3F	CC BF
00 03 40	03 40	CD 80
00 03 41	03 41	CD 81
.		
.		
.		
00 03 7E	03 7E	CD BE
00 03 7F	03 7F	CD BF
00 03 80	03 80	CE 80
00 03 81	03 81	CE 81
.		
.		
.		

Mappings, 4

UTF-32	UTF-16	UTF-8
00 03 BE	03 BE	CE BE
00 03 BF	03 BF	CE BF
00 03 C0	03 C0	CF 80
00 03 C1	03 C1	CF 81
.		
.		
.		
00 03 FE	03 FE	CF BE
00 03 FF	03 FF	CF BF
00 04 00	04 00	D0 80
00 04 01	04 01	D0 81
.		
.		
.		
00 04 3E	04 3E	D0 BE
00 04 3F	04 3F	D0 BF
00 04 40	04 40	D1 80
00 04 41	04 41	D1 81
.		
.		
.		
00 04 7E	04 7E	D1 BE
00 04 7F	04 7F	D1 BF
00 04 80	04 80	D2 80
00 04 81	04 81	D2 81
.		
.		
.		

Mappings, 5

UTF-32	UTF-16	UTF-8
00 04 BE	04 BE	D2 BE
00 04 BF	04 BF	D2 BF
00 04 C0	04 C0	D3 80
00 04 C1	04 C1	D3 81
.		
.		
.		
00 04 FE	04 FE	D3 BE
00 04 FF	04 FF	D3 BF
00 05 00	05 00	D4 80
00 05 01	05 01	D4 81
.		
.		
.		
00 05 3E	05 3E	D4 BE
00 05 3F	05 3F	D4 BF
00 05 40	05 40	D5 80
00 05 41	05 41	D5 81
.		
.		
.		
00 05 7E	05 7E	D5 BE
00 05 7F	05 7F	D5 BF
00 05 80	05 80	D6 80
00 05 81	05 81	D6 81
.		
.		
.		

Mappings, 6

UTF-32	UTF-16	UTF-8
00 05 BE	05 BE	D6 BE
00 05 BF	05 BF	D6 BF
00 05 C0	05 C0	D7 80
00 05 C1	05 C1	D7 81
.		
.		
.		
00 05 FE	05 FE	D7 BE
00 05 FF	05 FF	D7 BF
00 06 00	06 00	D8 80
00 06 01	06 01	D8 81
.		
.		
.		
00 06 3E	06 3E	D8 BE
00 06 3F	06 3F	D8 BF
00 06 40	06 40	D9 80
00 06 41	06 41	D9 81
.		
.		
.		
00 06 7E	06 7E	D9 BE
00 06 7F	06 7F	D9 BF
00 06 80	06 80	DA 80
00 06 81	06 81	DA 81
.		
.		
.		

Mappings, 7

UTF-32	UTF-16	UTF-8
00 06 BE	06 BE	DA BE
00 06 BF	06 BF	DA BF
00 06 C0	06 C0	DB 80
00 06 C1	06 C1	DB 81
.		
.		
.		
00 06 FE	06 FE	DB BE
00 06 FF	06 FF	DB BF
00 07 00	07 00	DC 80
00 07 01	07 01	DC 81
.		
.		
.		
00 07 3E	07 3E	DC BE
00 07 3F	07 3F	DC BF
00 07 40	07 40	DD 80
00 07 41	07 41	DD 81
.		
.		
.		
00 07 7E	07 7E	DD BE
00 07 7F	07 7F	DD BF
00 07 80	07 80	DE 80
00 07 81	07 81	DE 81
.		
.		
.		

Mappings, 8

UTF-32	UTF-16	UTF-8
00 07 BE	07 BE	DE BE
00 07 BF	07 BF	DE BF
00 07 C0	07 C0	DF 80
00 07 C1	07 C1	DF 81
.		
.		
.		
00 07 FE	07 FE	DF BE
00 07 FF	07 FF	DF BF
00 08 00	08 00	E0 A0 80
00 08 01	08 01	E0 A0 81
.		
.		
.		
00 08 3E	08 3E	E0 A0 BE
00 08 3F	08 3F	E0 A0 BF
00 08 40	08 40	E0 A1 80
00 08 41	08 41	E0 A1 81
.		
.		
.		
00 08 7E	08 7E	E0 A1 BE
00 08 7F	08 7F	E0 A1 BF
00 08 80	08 80	E0 A2 80
00 08 81	08 81	E0 A2 81
.		
.		
.		

Mappings, 9

UTF-32	UTF-16	UTF-8
00 08 BE	08 BE	E0 A2 BE
00 08 BF	08 BF	E0 A2 BF
00 08 C0	08 C0	E0 A3 80
00 08 C1	08 C1	E0 A3 81
.		
.		
.		
00 08 FE	08 FE	E0 A3 BE
00 08 FF	08 FF	E0 A3 BF
00 09 00	09 00	E0 A4 80
00 09 01	09 01	E0 A4 81
.		
.		
.		
00 09 3E	09 3E	E0 A4 BE
00 09 3F	09 3F	E0 A4 BF
00 09 40	09 40	E0 A5 80
00 09 41	09 41	E0 A5 81
.		
.		
.		

-- a big jump here, since the pattern is established --

00 0F FE	0F 7E	E0 BF BE
00 0F FF	0F 7F	E0 BF BF
00 10 00	10 00	E1 80 80
00 10 01	10 01	E1 80 81
.		
.		
.		

Mappings, 10

UTF-32	UTF-16	UTF-8
-- another big jump here, since the pattern is established --		
00 1B FE	1B 7E	E1 AF BE
00 1B FF	1B 7F	E1 AF BF
00 1C 00	1C 00	E1 B0 80
00 1C 01	1C 01	E1 B0 81
.		
.		
.		
-- another big jump here, since the pattern is established --		
00 1F FE	1F 7E	E1 BF BE
00 1F FF	1F 7F	E1 BF BF
00 20 00	20 00	E2 80 80
00 20 01	20 01	E2 80 81
.		
.		
.		
-- another big jump here, since the pattern is established --		
00 2F FE	2F 7E	E2 BF BE
00 2F FF	2F 7F	E2 BF BF
00 30 00	30 00	E3 80 80
00 30 01	30 01	E3 80 81
.		
.		
.		
-- another big jump here, since the pattern is established --		
00 3F FE	3F 7E	E3 BF BE
00 3F FF	3F 7F	E3 BF BF
00 40 00	40 00	E4 80 80
00 40 01	40 01	E4 80 81
.		
.		
.		

Mappings, 11

UTF-32

UTF-16

UTF-8

-- another big jump here, since the pattern is established --

00	4F	FE	4F	7E	E4	BF	BE
00	4F	FF	4F	7F	E4	BF	BF
00	50	00	50	00	E5	80	80
00	50	01	50	01	E1580	81	

.
. .

-- another big jump here, since the pattern is established --

00	D7	FE	D7	FE	ED	9F	BE
00	D7	FF	D7	FF	ED	9F	BF

-- at this point, we have reached the place where
surrogate characters occur; a surrogate character
by itself is not a valid Unicode character;
we pick up again, after the surrogate points:

00	F9	00	F9	00	EF	A4	80
00	F9	01	F9	01	EF	A4	81

.
. .

00	FF	FE	FF	FE	EF	BF	BE
00	FF	FF	FF	FE	EF	BF	BF

-- the next code point, x'010000', and all code points
after this, will require pairs of surrogate characters
for the UTF-16 vlaues ...

Mappings, 12

UTF-32	UTF-16	UTF-8
01 00 00	D8 00 DC 00	F0 90 80 80
01 00 01	D8 00 DC 01	F0 90 80 81
.		
.		
.		
01 01 FE	D8 00 DD FE	F0 90 87 BE
01 01 FF	D8 00 DD FF	F0 90 87 BF
01 02 00	D8 00 DE 00	F0 90 88 80
01 02 01	D8 00 DE 01	F0 90 88 81
.		
.		
.		
01 02 FE	D8 00 DE FE	F0 90 88 BE
01 02 FF	D8 00 DE FF	F0 90 88 BF

-- another big jump here, since the pattern is established --

01 E0 00	D8 01 DC 00	F0 90 90 80
01 E0 01	D8 01 DC 01	F0 90 90 81

-- another big jump here, since the pattern is established --

01 20 00	D8 40 DC 00	F0 90 90 80
01 20 01	D8 41 DC 01	F0 90 90 81

Mappings, 12

UTF-32

UTF-16

UTF-8

**-- a final big jump here, up to the end of
assigned characters**

0E 00 00	DB 40 DC 00	F3 A0 80 80
0E 00 01	DB 40 DC 01	F3 A0 80 81
.		
.		
.		
0E 01 EE	DB 40 DD EE	F3 A0 87 AE
0E 01 EF	DB 40 DD EF	F3 A0 87 AF